

Markov Chain Monte Carlo

Introduction, Comparison & Analysis

Tzu-Heng Lin, 2014011054, W42

Department of Electronic Engineering, Tsinghua University, Beijing, China
lzhbrian@gmail.com

ABSTRACT

Markov Chain Monte Carlo (MCMC) is a technique to make an estimation of a statistic by simulation in a complex model. Restricted Boltzmann Machine(RBM) is a crucial model in the field of Machine Learning. However, training a large RBM model will include intractable computation of the partition functions, i.e. $Z(\theta)$. This problem has aroused interest in the work of estimation using a MCMC methods. In this paper, we first conduct Metropolis-Hastings Algorithm, one of the most prevalent sampling methods, and analyze its correctness & performance, along with the choice of the accepting rate. We then implement three algorithms: TAP, AIS, RTS, to estimate partition functions of an RBM model. Our work not only give an introduction about the available algorithms, but systematically compare the performance & difference between them. We seek to provide an overall view in the field of MCMC.

1. INTRODUCTION

Markov Chain A Markov Chain is a special stochastic process in which the current state only depends on its previous state, we call this property the Markov property or memorylessness. i.e.

$$P(X_{t+1} = x | X_t, X_{t-1}, \dots) = P(X_{t+1} = x | X_t) \quad (1)$$

Given a Markov Chain, if a vector π , has the property:

$$\pi = \pi P \quad (2)$$

then we call π the stationary distribution, it denotes the final state distribution of the stochastic process in a Markov Chain.

Markov Chain has a wide application in many fields in the real world, such as social science[1], economics & finance[7] and of course, computer science[16], etc.

Markov Chain Monte Carlo If we are given a probability distribution $p(x)$, it would be a great thing if we could generate some samples of it by a simple method, so here comes the Markov Chain Monte Carlo(MCMC) method. If we could construct a Markov Chain which its stationary distribution π just equals to $p(x)$, then we could use this Markov Chain to sample from this distribution $p(x)$. And this is the main idea of MCMC.

In this paper, we implement the Metropolis-Hastings Algorithm[14, 8], which is one of the most widely used sampling method. We also deal with the accepting rate, which I will introduce later, for previous work[18] have shown that the accepting rate may influence the result of the experiment

and there is theoretical support in choosing an optimal accepting rate.

Restricted Boltzmann Machine A Restricted Boltzmann Machine (RBM)[13] is a significant work bringing hypothesis in statistical physics to computer science. By stacking several layers of RBM, we will get a fundamental model, Deep Belief Network[9], in the field of Deep Learning, which is nowadays the hottest class of algorithms used in Machine Learning.

Estimating Partition functions In the process of training an RBM, however, will include intractable computation of the partition function. When the model grows large, the complexity of this work will be incompletable. The good news is, several researches have shown that there are ways to avoid this by using an MCMC approach instead, to estimate it.

In this paper, we implement three prevalent MCMC methods of estimating a partition function, Thouless-Anderson-Palmer Sampling(TAP)[5], Annealed Importance Sampling(AIS)[15, 19], Rao-Blackwellized Tempered Sampling(RTS)[3], respectively, and give an overall comparison on the theory & performance between them.

2. METROPOLIS-HASTINGS

In the Introduction, we have shown the main idea of an MCMC sampling method. In this section, we will introduce the Metropolis Hastings Algorithm and conduct an experiment.

2.1 Algorithm¹

2.1.1 Detailed Balance Condition

Before stepping further into the MH Algorithm, We would first introduce a theorem called the Detailed Balance Condition.

In the introduction, we said that we want to construct a Markov Chain which its stationary distribution $\pi(x)$ just equals to the required probability distribution $p(x)$.

At first, a theorem is needed.

THEOREM 2.1 (DETAILED BALANCE CONDITION). *Given a non periodic Markov Chain, if*

$$\pi(i)P_{ij} = \pi(j)P_{ji} \quad \text{for all } i, j \quad (3)$$

then $\pi(x)$ is the stationary distribution of this Markov Chain.

¹Available at https://github.com/lzhbrian/MCMC/blob/master/metropolis_hastings/metropolis_hasting.R in R[17]

2.1.2 MCMC sampling method

Suppose we already have a transition matrix Q for a Markov Chain, $q(i, j)$ denote the probability of transition from state i to state j . For the general case,

$$p(i)q(i, j) \neq p(j)q(j, i)$$

That is to say, we do not have the detailed balance condition (Theorem 2.1) So we introduce an $\alpha(i, j)$ s.t.

$$p(i)q(i, j)\alpha(i, j) = p(j)q(j, i)\alpha(j, i) \quad (4)$$

By symmetrical characteristic, we choose:

$$\alpha(i, j) = p(j)q(j, i) \quad \alpha(j, i) = p(i)q(i, j) \quad (5)$$

So the new Markov Chain Q' would have the property of which its stationary distribution is $p(x)$

$$p(i) \underbrace{q(i, j)\alpha(i, j)}_{Q'(i, j)} = p(j) \underbrace{q(j, i)\alpha(j, i)}_{Q'(j, i)} \quad (6)$$

We call the $\alpha(i, j)$ we introduced, accepting ratio. It means that, in the original Markov Chain Q , when state i transits to state j with a probability of $q(i, j)$, we accept this transition with a probability of $\alpha(i, j)$

Now, we have derived the MCMC sampling method.

2.1.3 Metropolis-Hastings Algorithm

The MCMC sampling method is a marvellous work. However, it has a critical drawback that if $\alpha(i, j)$ & $\alpha(j, i)$ are too small, we would seldom accept the transition.

A solution is that we multiply both $\alpha(i, j)$ & $\alpha(j, i)$ with a constant to make sure that the larger one between them equals 1. By doing so, we change the accepting ratio to

$$\alpha(i, j) = \min \left\{ \frac{p(j)q(j, i)}{p(i)q(i, j)}, 1 \right\} \quad (7)$$

and now, we get Metropolis-Hastings Algorithm[8].

I would like to further introduce one more concept called accepting rate(not accepting ratio), which denotes the statistic ratio of accepting the transition. i.e. If we request 10 transition and we accept 8 times, then the accepting rate would be 0.8. This concept is crucial when we are dealing with a continual Markov Chain to use the MH algorithm.

2.1.4 Symmetric Case

In a Markov Chain whose transition matrix is symmetric, we have

$$q(i, j) = q(j, i)$$

so the accepting ratio could be simplified to

$$\alpha(i, j) = \min \left\{ \frac{p(j)}{p(i)}, 1 \right\} \quad (8)$$

which is also known as the Metropolis Algorithm[14].

2.1.5 Continual Case

In a continual Markov Chain, such as the experiment we are going to do in the next subsection, we have a vague definition of transition matrix Q . So we introduce a concept called the proposal jump size, $sd.T$.

The method we get x_{k+1} from x_k is to add a sampled point of a normal distribution with a variance of the jump size and $\mu = 0$. For a two dimension example, we have:

$$x_{k+1} = x_k + sd.T \begin{pmatrix} norm_1 \\ norm_2 \end{pmatrix} \quad (9)$$




Figure 1: Sampling result of 5,000 points correlation = 0.50009 , set sd.T = 3.0

2.2 Sampling Experiment

For our experiment, we use an example of a bivariate Normal distribution, with

$$\mu = \begin{pmatrix} 5 \\ 10 \end{pmatrix}, \Sigma = \begin{pmatrix} 1 & 1 \\ 1 & 4 \end{pmatrix}$$

By theoretical computation, we can easily compute the the pearson correlation between the two dimensional value is 0.5.

$$\rho = 0.5$$

We then generate 10,000 samples using the MH algorithm and take the second half (i.e. the last 5,000 points), setting the standard deviation of proposal to 3.0. We can see from the result (Figure 1) that we have derived 5,000 sampled points whose pearson correlation value $\rho = 0.50009$, which matches the theoretical value.

2.3 Performance Analysis

2.3.1 Choice of proposal jump size

MH algorithm is an effective MCMC method for many diverse problems. However, its efficiency depends crucially on the selection of the proposal density. With the proposal jump size being small, the accepting rate would be very low and eventually stick to only one point(eg. the initial point); When the proposal jump size is too big, the accepting rate would be too high.

Roberts et al. have shown in previous work[18] that the optimal accepting rate of the MH algorithm should approximately be at 0.234 for the case of an N-dimensional Gaussian target distribution. We test the accepting rates in different proposal jump size(Figure 2) and find that the optimal value should be at approximately 3.0 to acquire a model with accepting rate being close to 0.234. That is the reason why we choose 3.0 as our proposal jump size.

2.3.2 Efficiency

Due to the limit of the accepting rate, for a high dimensional condition, using the MH sampling methods may spend more time in traverse all of the possible states, which could sometimes be less satisfying. Thus, many would switch to Gibbs Sampling Algorithm.

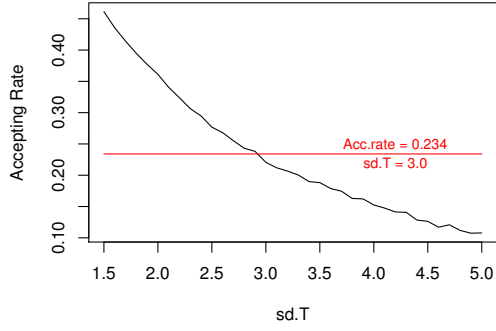


Figure 2: Accepting rate on different proposal jump size.

2.4 Gibbs Sampling

Gibbs Sampling is a special case of Metropolis Hastings Algorithm, by letting the accepting rate = 1, we will get a Gibbs Sampler. As the length & time limit, we will not specify more here.

3. PARTITION FUNCTION ESTIMATION

3.1 Restricted Boltzmann Machine

Co-invented and enhanced largely[9] by Geoff Hinton, a Restricted Boltzmann Machine(RBM)[13] is a model which brings the idea of a physics concept to the field of computer science.

3.1.1 Introduction

An RBM is a two-layer undirected model(Figure 3). The first layer of the RBM is called visible layer, and the second is called the hidden layer. In the model, every visible units are connected to all hidden units and vice versa. For every given value of visible layer \mathbf{v} & hidden layer \mathbf{h} , we can define an energy of this state.

$$E(\mathbf{v}, \mathbf{h}; \theta) = -\mathbf{v}^T \mathbf{W} \mathbf{h} - \mathbf{b}^T \mathbf{v} - \mathbf{a}^T \mathbf{h} \quad (10)$$

where $\theta = \{W, \mathbf{b}, \mathbf{a}\}$ are the model configurations. W_{ij} represents the weight between visible unit v_i and hidden unit h_j . \mathbf{b} & \mathbf{a} are biases for visible and hidden layer, respectively.

3.1.2 Training an RBM

On training an RBM, we want our RBM model to have a lowest scale of energy. By doing so, we have to calculate the joint distribution over the visible and hidden units, which is defined by:

$$p(\mathbf{v}, \mathbf{h}; \theta) = \frac{e^{-E(\mathbf{v}, \mathbf{h}; \theta)}}{Z(\theta)} \quad (11)$$

where

$$Z(\theta) = \sum_{\mathbf{v}} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h}; \theta)} \quad (12)$$

is the partition function.

However, calculating partition functions has always been an intractable work since we have to traverse all the possible

state of \mathbf{v} & \mathbf{h} . When the model grows large, this process will be very time & resources consuming and thus become unrealistic for the real practice.

So, we have to introduce methods to estimate the partition functions instead of just calculating it in brute force. Although some deviation may include in the estimation, but the efficiency along with them make them preferable. In fact, studies have shown that only few deviation is included that we could just ignore it since it does petty influence on our training.

In the next subsection, we will discuss about three methods available, which each have their pros and cons in doing this complex estimation.

3.2 Algorithms

3.2.1 Thouless-Anderson-Palmer Sampling²

Thouless-Anderson-Palmer Sampling(TAP)[5] is a very efficient and easy-to-practice iterative procedure based on an improved mean field method from statistical physics called Thouless-Anderson-Palmer approach.

The main idea of this method is to iteratively compute the magnetization vector m^v, m^h , and then input the values into the Legendre transform of the free energy $F = \log(Z(\theta))$ to compute it.

The Legendre transform of F to the second order is:

$$\begin{aligned} \Gamma(\mathbf{m}^v, \mathbf{m}^h) \approx & -S(\mathbf{m}^v, \mathbf{m}^h) - \sum_i a_i m_i^v - \sum_j b_j m_j^h \\ & - \sum_{i,j} \left(W_{i,j} m_i^v m_j^h \right. \\ & \left. - 0.5 W_{i,j} (m_i^v - (m_i^v)^2) (m_j^h - (m_j^h)^2) \right) \end{aligned} \quad (13)$$

where $S(\mathbf{m}^v, \mathbf{m}^h)$ indicates the entropy:

$$\begin{aligned} S(\mathbf{m}^v, \mathbf{m}^h) = & - \sum_i \left(m_i^v \log m_i^v + (1 - m_i^v) \log(1 - m_i^v) \right) \\ & - \sum_j \left(m_j^h \log m_j^h + (1 - m_j^h) \log(1 - m_j^h) \right) \end{aligned} \quad (14)$$

The pseudo code of this algorithm is shown below:

3.2.2 Annealed Importance Sampling³

Algorithm Annealed Importance Sampling(AIS)[15, 19] is probably one of the most preferable estimating methods available.

Previous work [12] have shown that if P_A and P_B in the SIS method is not close enough, the estimator would be very poor.

Based on SIS, the main idea of this algorithm is to gradually alter the value from an known Z_A to our required Z_B (or Z_K), by the following identity:

$$\frac{Z_K}{Z_0} = \frac{Z_1}{Z_0} \frac{Z_2}{Z_1} \dots \frac{Z_K}{Z_{K-1}} \quad (15)$$

²Available at <https://github.com/lzhhbrian/MCMC/blob/master/rbm/TAP.m> in Matlab

³Available at <https://github.com/lzhhbrian/MCMC/blob/master/rbm/AIS.m> in Matlab




Figure 3: A Restricted Boltzmann Machine




Figure 4: The transition process from x_k to x_{k+1} which leaves $P_k(\mathbf{v})$ invariant.

where

$$\frac{Z_K}{Z_{k+1}} = \frac{1}{M} \sum_{i=1}^M \frac{P_{k+1}^*(\mathbf{x}^{(i)})}{P_k^*(\mathbf{x}^{(i)})} \text{ where } \mathbf{x}^{(i)} \sim P_k \quad (16)$$

in which we can get x_{k+1} from:

$$\begin{aligned} p(h_j^A = 1 | \mathbf{v}) &= \text{sigmoid} \left((1 - \beta_k) \left(\sum_i W_{ij}^A v_i + a_j^A \right) \right) \\ p(h_j^B = 1 | \mathbf{v}) &= \text{sigmoid} \left(\beta_k \left(\sum_i W_{ij}^B v_i + a_j^B \right) \right) \\ p(v'_i = 1 | \mathbf{h}) &= \text{sigmoid} \left((1 - \beta_k) \left(\sum_j W_{ij}^A h_j^A + b_i^A \right) \right. \\ &\quad \left. + \beta_k \left(\sum_j W_{ij}^B h_j^B + b_i^B \right) \right) \end{aligned} \quad (17)$$

this procedure is shown in Figure 4.

Note that model A indicates an initial model which we can easily compute all its configurations. Commonly, we choose an RBM model with $\theta = \{0, 0, 0\}$

β in the above equations is defined by users as a set of inverse temperatures $\{0 = \beta_1 < \beta_2 < \dots < \beta_K = 1\}$, which can define a sequence of

$$P_k(\mathbf{x}) \propto P_A^*(\mathbf{x})^{1-\beta_k} P_B^*(\mathbf{x})^{\beta_k} \quad (18)$$

where

$$P_k^*(\mathbf{v}) = \sum_{\mathbf{h}^A \mathbf{h}^B} e^{(1-\beta_k)E(\mathbf{v}, \mathbf{h}^A; \theta_A) + \beta_k E(\mathbf{v}, \mathbf{h}^B; \theta_B)} \quad (19)$$

Initialize Z_A with dataset In [19], Ruslan et al. also notice a method to make Z_A near Z_B . As the length & time limit, we will not specify it here.

The pseudo code of this algorithm is shown below:

3.2.3 Rao-Blackwellized Tempered Sampling⁴

Algorithm Similar to AIS, Rao-Blackwellized Tempered (RTS)[3] Sampling also has a set of inverse temperatures $\{0 = \beta_1 < \beta_2 < \dots < \beta_K = 1\}$, which can define a sequence of

$$f_k(\mathbf{x}) \propto f(\mathbf{x})^{\beta_k} p_1(\mathbf{x})^{1-\beta_k} \quad (20)$$

Different from AIS, we do not traverse β . Instead we sample a β^* every loop, from the β set with the distribution $(\beta|x)$.

Subsequently, we sample from x_k to x_{k+1} by the probability of $q(x|\beta^*)$ just like what we did in AIS, shown in Figure 4. However, what also different from AIS is that, we have to iterate from x_k to x_{k+1} many times (i.e. 50 times in [3]) for the sake of getting a better x_{k+1} .

At the last of every loop, we update the lower variance estimator \hat{c} by

$$\hat{c}_k = \hat{c}_k + \frac{1}{N} q(\beta_k | x) \quad (21)$$

Finally, we get Z_k by

$$\hat{Z}_k^{RTS} = \hat{Z}_k \frac{r_1 \hat{c}_k}{r_k \hat{c}_1}, \quad k = 2, \dots, K \quad (22)$$

in which what we do care is $Z_B \approx \hat{Z}_K^{RTS}$.

The posterior distribution $q(\beta_k | x)$ in the above equations is defined by:

$$q(\beta_k | x) = \frac{f_k(x) r_k / \hat{Z}_k}{\sum_{k'=1}^K f_{k'}(x) r_{k'} / \hat{Z}_{k'}} \quad (23)$$

Practice In the paper [3], Carlson et al. note an initializing method to initialize Z_k , whose procedure is just like the above process. The only difference is that they sampled β_k by uniform distribution in every loop, not by the distribution $(\beta|x)$. They claim that after doing such initializing work, then we conduct the algorithms above would acquire a better result.

⁴Available at <https://github.com/lzhbrian/MCMC/blob/master/rbm/RTS.m> in Matlab

In our real practice, we directly use the initializing method mentioned above by selecting β_k with a uniform distribution in every loop. We also initialize the value of Z_A by the method we have mentioned in the AIS section using the dataset. And we have found that the result is already satisfying, there is no need to conduct more loops with β_k sampled by $(\beta|x)$.

Also, we found that we have to conduct the procedure above for several times s.t. we can acquire our desired partition function value.(i.e. We did it for 100 times, that is to say we update \mathbf{Z} for 100 times).

The pseudo code of this algorithm is shown below:

3.2.4 Other method

There are many other methods which can also estimate the partition functions. Such as Self-adjusted mixture sampling(SAMS)[20] proposed a method to estimate multiple partition functions together to improve the efficiency. As the length & time limit, we only implement 3 methods here in this paper.

3.3 Estimating Results

If we don't consider about the time & resources, we can get good results for all three methods(Figure x).

3.4 Performance Analysis

3.4.1 Practice

Complexity For horizontal comparison on complexity, we let the correctness & stability to be approximately the same (i.e. the correctness the TAP could achieve in 2 seconds), and compare the run time for three algorithms.

We have conducted 20 experiments for the acquisition of $Z(\theta)$.

From our practice, we have found that xxx gives the most agreeable result while TAP completely underestimate the value.

Correctness & Stability For horizontal comparison on the correctness & stability, we allow the same run time for three algorithms to compute (i.e. 10 seconds per experiment). We have conducted 20 experiments for the acquisition of $Z(\theta)$.

From our practice, we have found that xxx gives the most agreeable result while TAP completely underestimate the value.

We also notice that the variance of ... is the smallest which indicates it has the best stability.

3.4.2 Theoretical

RTS From a theoretical perspective, we have proven that the bias & the variance of the RTS method are to be:

$$E[\log \hat{Z}_k^{RTS}] - E[Z_k] \approx \frac{1}{2} \left[\frac{\sigma_1^2}{\hat{c}_1^2} - \frac{\sigma_k^2}{\hat{c}_k^2} \right] \quad (24)$$

$$Var[\log \hat{Z}_k^{RTS}] \approx \frac{\sigma_1^2}{\hat{c}_1^2} + \frac{\sigma_k^2}{\hat{c}_k^2} - \frac{2\sigma_{1k}}{\hat{c}_k \hat{c}_1} \quad (25)$$

where $\sigma_k^2 = Var[\hat{c}_k]$ and $\sigma_{1k} = Cov[\hat{c}_1, \hat{c}_k]$

This has shown that the bias of RTS has no definite sign.

AIS However, in AIS, Neal and Jarzynski et al.[15, 11] have shown that if we want the result to be unbiased, we would have to let $M = 1$ in the iteration, which by doing so

have lost the advantage of AIS. On the other hand, if $M > 1$, we would have a negative bias due to Jensen Inequality.

TAP Although TAP shows the best efficiency, its results are the most disappointing. Apparently TAP has underestimate the value of the partition function.

We did not analyze deep on the reason why it failed to perform a satisfying result, but our intuition tell that maybe it is because the Legendre transform. In our practice, we only took the Legendre transform to the 2nd order, which might result in the underestimation.

4. RELATED WORK

In this paper, we do not dig very deep into a specific field of the algorithms, but provide an overall view of the MCMC method. There are brilliant seniors who have done many marvellous works in different specific field.

Complexity Hubbard[10] is probably the earliest researcher managing to reduce the complexity of a partition function. Beskos et al.[2] made a great work in analysing the complexity of the Metropolis-Hastings Algorithm. Gilks[6] did a significant job in introducing and expansion of the MCMC method.

Theoretical Explanation Chib et al.[4] gave an clear introduction and deeper explanation to the Metropolis-Hastings Algorithm. Jarzynski[11] systematically analyzed the bias between the estimation and the real value of the partition functions in a sampling method.

Due to my limited knowledge, there might be some mistakes and flaws in this paper, please don't hesitate to contact and correct me.

5. CONCLUSION

In this paper, we discuss about the Markov Chain Monte Carlo method which are now undoubtedly one of the most important sampling methods.

We comprehensively introduce the concept of Metropolis-Hastings Algorithm and conduct an experiment to verify its correctness. We also make some analysis about how accepting rate would interfere the sampling result.

We systematically compare three methods of partition function estimation which are crucial works in training a Restricted Boltzmann Machine or a Deep Belief Network.

As future work, we would like to join more methods to the comparison and if could, propose some improvement to the algorithms available.

6. ACKNOWLEDGEMENT

I would like to thank Yubo Chen, Liren Yu, Yuanxin Zhang, XueChao Wang, Changran Hu, for the discussion with me on the algorithms. Without them, I wouldn't have the possibility to accomplish this work in such a short time. This paper is a project of Stochastic Process Course in Tsinghua University, taught by Prof. Zhijian Ou.

7. REFERENCES

- [1] D. Acemoglu, G. Egorov, and K. Sonin. Political model of social evolution. *Proceedings of the National Academy of Sciences*, 108(Supplement 4):21292–21296, 2011.
- [2] A. Beskos and A. Stuart. Computational complexity of metropolis-hastings methods in high dimensions. In *Monte Carlo and Quasi-Monte Carlo Methods 2008*, pages 61–71. Springer, 2009.

- [3] D. Carlson, P. Stinson, A. Pakman, and L. Paninski. Partition functions from rao-blackwellized tempered sampling. *arXiv preprint arXiv:1603.01912*, 2016.
- [4] S. Chib and E. Greenberg. Understanding the metropolis-hastings algorithm. *The american statistician*, 49(4):327–335, 1995.
- [5] M. Gabrié, E. W. Tramel, and F. Krzakala. Training restricted boltzmann machine via the thouless-anderson-palmer free energy. In *Advances in Neural Information Processing Systems*, pages 640–648, 2015.
- [6] W. R. Gilks. *Markov chain monte carlo*. Wiley Online Library, 2005.
- [7] J. D. Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica: Journal of the Econometric Society*, pages 357–384, 1989.
- [8] W. K. Hastings. Monte carlo sampling methods using markov chains and their applications. *Biometrika*, 57(1):97–109, 1970.
- [9] G. E. Hinton and R. R. Salakhutdinov. Reducing the dimensionality of data with neural networks. *Science*, 313(5786):504–507, 2006.
- [10] J. Hubbard. Calculation of partition functions. *Physical Review Letters*, 3(2):77, 1959.
- [11] C. Jarzynski. Nonequilibrium equality for free energy differences. *Physical Review Letters*, 78(14):2690, 1997.
- [12] D. J. MacKay. *Information theory, inference and learning algorithms*. Cambridge university press, 2003.
- [13] J. L. McClelland, D. E. Rumelhart, P. R. Group, et al. Parallel distributed processing, 1987.
- [14] N. Metropolis, A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller, and E. Teller. Equation of state calculations by fast computing machines. *The journal of chemical physics*, 21(6):1087–1092, 1953.
- [15] R. M. Neal. Annealed importance sampling. *Statistics and Computing*, 11(2):125–139, 2001.
- [16] L. Page, S. Brin, R. Motwani, and T. Winograd. The pagerank citation ranking: bringing order to the web. 1999.
- [17] R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2016.
- [18] G. O. Roberts, A. Gelman, W. R. Gilks, et al. Weak convergence and optimal scaling of random walk metropolis algorithms. *The annals of applied probability*, 7(1):110–120, 1997.
- [19] R. Salakhutdinov. *Learning deep generative models*. PhD thesis, University of Toronto, 2009.
- [20] Z. Tan. Optimally adjusted mixture sampling and locally weighted histogram analysis. *Journal of Computational and Graphical Statistics*, (just-accepted), 2015.